

Know Your (Java) App

Collect metrics with Micrometer

Erin Schnabel @ebullientworks

What makes a system (or service) observable?

Workload routing
System health

Health Checks



Service is ready
Service is not a zombie

Statistics & trends
Analytics

Metrics



How many times was
method x called?

Service-centric
problem determination

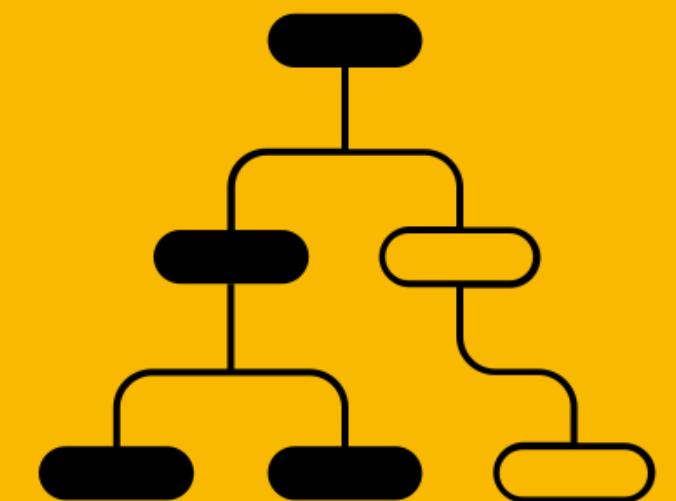
Log Entries



What happened when
method x was called

Context + relationships
for end-to-end analysis

Distributed Trace



Method x was called

Metrics: Time series data

How does data change *over time*?

- Time is the x-axis
 - Data collected at regular intervals
 - Each measured value appended to the end of a series
- Used for **trend analysis, anomaly detection, alerting**
Not for debugging, diagnostics, or root cause analysis

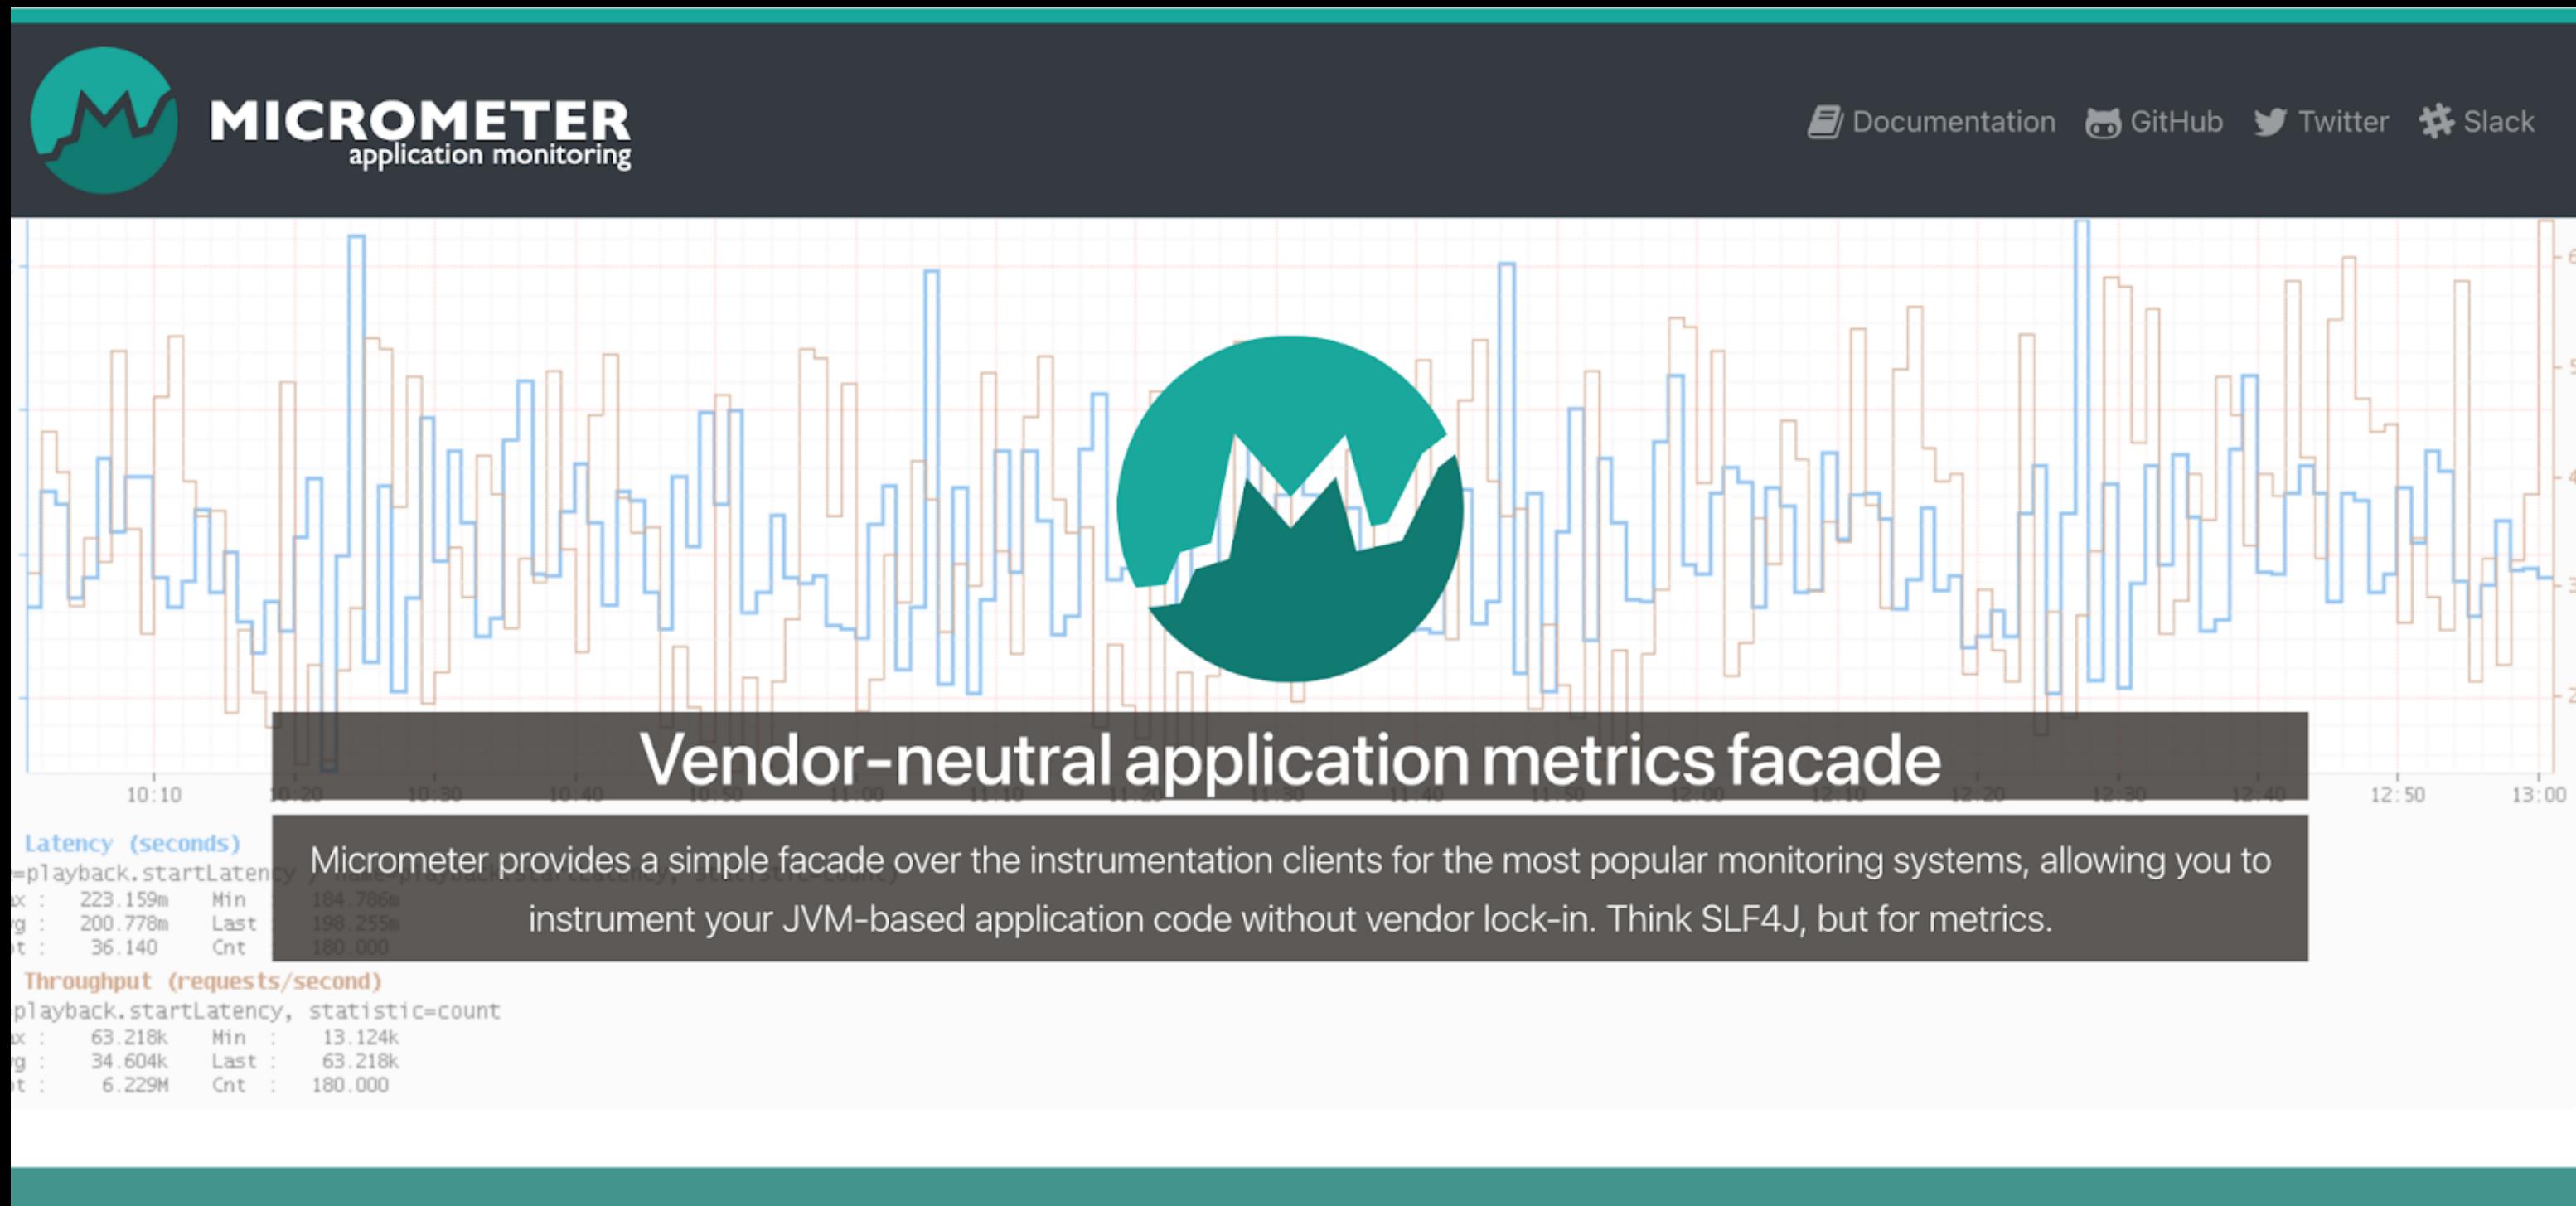
Metrics: Dimensions

Filtered>Selective aggregation

```
http_server_requests_seconds_count{  
    env="test",  
    method="GET",  
    outcome="SUCCESS",  
    status="200",  
    uri="/example/prime/{number}" ,} 5.0
```

Cardinality: Each unique combination (name & all tag/values) is its own time series

What is Micrometer?



"Screw gauge" micrometer

<https://micrometer.io> -- Open Source metrics library

Developer API

- Obtain a **MeterRegistry**
- Create a **Meter** (name + tag=value combination)

Gauge	FunctionCounter
Counter	FunctionTimer, LongTaskTimer
Timer	
DistributionSummary	

Counter: increase only → `.increment()`

1. `registry.counter("example.prime.number", "type", "prime");`
2. `Counter.builder("counter")
 .baseUnit("beans") // optional
 .description("a description") // optional
 .tags("region", "test") // optional
 .register(registry);`
3. `@Counted(value = "metric.all", extraTags = { "region", "test" }) **`

Timer → .record(duration)

1. `registry.timer("my.timer", "region", "test");`
2. `Timer.builder("my.timer")
 .description("description ") // optional
 .tags("region", "test") // optional
 .register(registry);`
3. `@Timed(value = "call", extraTags = {"region", "test"}) **`

Working with Timers

1. `timer.record(() → noReturnValue());`
2. `timer.recordCallable(() → returnValue());`
3. `Runnable r = timer.wrap(() → noReturnValue());`
4. `Callable c = timer.wrap(() → returnValue());`
5. `Sample s = Timer.start(registry); doStuff; s.stop(timer);`

Gauge: increase/decrease, observed value

```
1. List<String> list = registry.gauge("list.size", Tags.of( ... ),  
    new ArrayList<>(), List::size);  
  
2. Gauge.builder("jvm.threads.peak", threadBean, ThreadMXBean::getPeakThreadCount)  
    .tags("region", "test")  
    .description("The peak live thread count ...")  
    .baseUnit(BaseUnits.THREADS)  
    .register(registry);
```

FunctionCounter: gauge-like, counter function

```
1. registry.more().counter( ... )  
  
2. FunctionCounter.builder("hibernate.transactions", statistics,  
   s → s.getTransactionCount() - s.getSuccessfulTransactionCount())  
   .tags("entityManagerFactory", sessionFactoryName)  
   .tags("result", "failure")  
   .description("The number of transactions we know to have failed")  
   .register(registry);
```

FunctionTimer: gauge-like, counter & duration function

```
1. registry.more().timer( ... );  
  
2. FunctionTimer.builder("cache.gets.latency", cache,  
    c → c.getLocalMapStats().getGetOperationCount(),  
    c → c.getLocalMapStats().getTotalGetLatency(),  
    TimeUnit.NANOSECONDS)  
    .tags("name", cache.getName())  
    .description("Cache gets")  
    .register(registry);
```

LongTaskTimer (active count, longest active task)

```
1. registry.more().longTaskTimer("long.task", "region", "test");
```

```
2. LongTaskTimer.builder("long.task")
   .description("a description") // optional
   .tags("region", "test") // optional
   .register(registry);
```

```
3. @Timed(value = "long.task", extraTags = {"extra", "tag"}, longTask = true) **
```

Distribution Summary → .record(value), not time

```
1. registry.summary("name", "region", "test")  
  
2. DistributionSummary.builder("response.size")  
   .description("a description") // optional  
   .baseUnit("bytes")          // optional  
   .tags("region", "test")     // optional  
   .register(registry);
```

Library Maintainer API

- Provide a **MeterBinder** (optional dependency)
 - Register library-specific metrics
- Quarkus / Spring discover **MeterBinder** implementations
 - binder.bindTo(registry)

cache
commonspool2
db
grpc
http
httpcomponents
hystrix
jetty
jpa
jvm
kafka
logging
mongodb
okhttp3
system
tomcat

Configuration and Filtering

- Late-binding configuration provided by **MeterFilter**
 - Sourced from user code or shared library
- Accept/Deny filters: fine-grained control over gathered metrics
- Transform: rename metrics, add/ignore/replace/rename tags
- Customize Timer / Histogram behavior
 - Percentiles, quantiles, Service Level Agreement(SLA) boundaries

Monster Combat

<https://github.com/ebullient/monster-combat>

 **Erin Schnabel** @ebullientworks · Nov 17, 2020

Normal 8%

Never ever would I have guessed I would write something like this. Damn kids. ;)

Exploring Application Metrics with Dungeons & Dragons



Exploring Application Metrics with Dungeons & Dragons
Erin Schnabel explores what the popular role-playing game Dungeons & Dragons has in common with application metrics.

fosslife.org

1 3 8

...

Roll for initiative

Monsters in combat: exploring application metrics with D&D

Nov 9, 2020 ·  **Erin Schnabel**

#metrics #quarkus



What does one of the most popular pen-and-paper role-playing games have in common with application metrics and frameworks like Quarkus or Micrometer? This epic article takes you on an adventure full of monster battles, graphs and metrics. And now everyone grab a d20 and roll to “Read”!

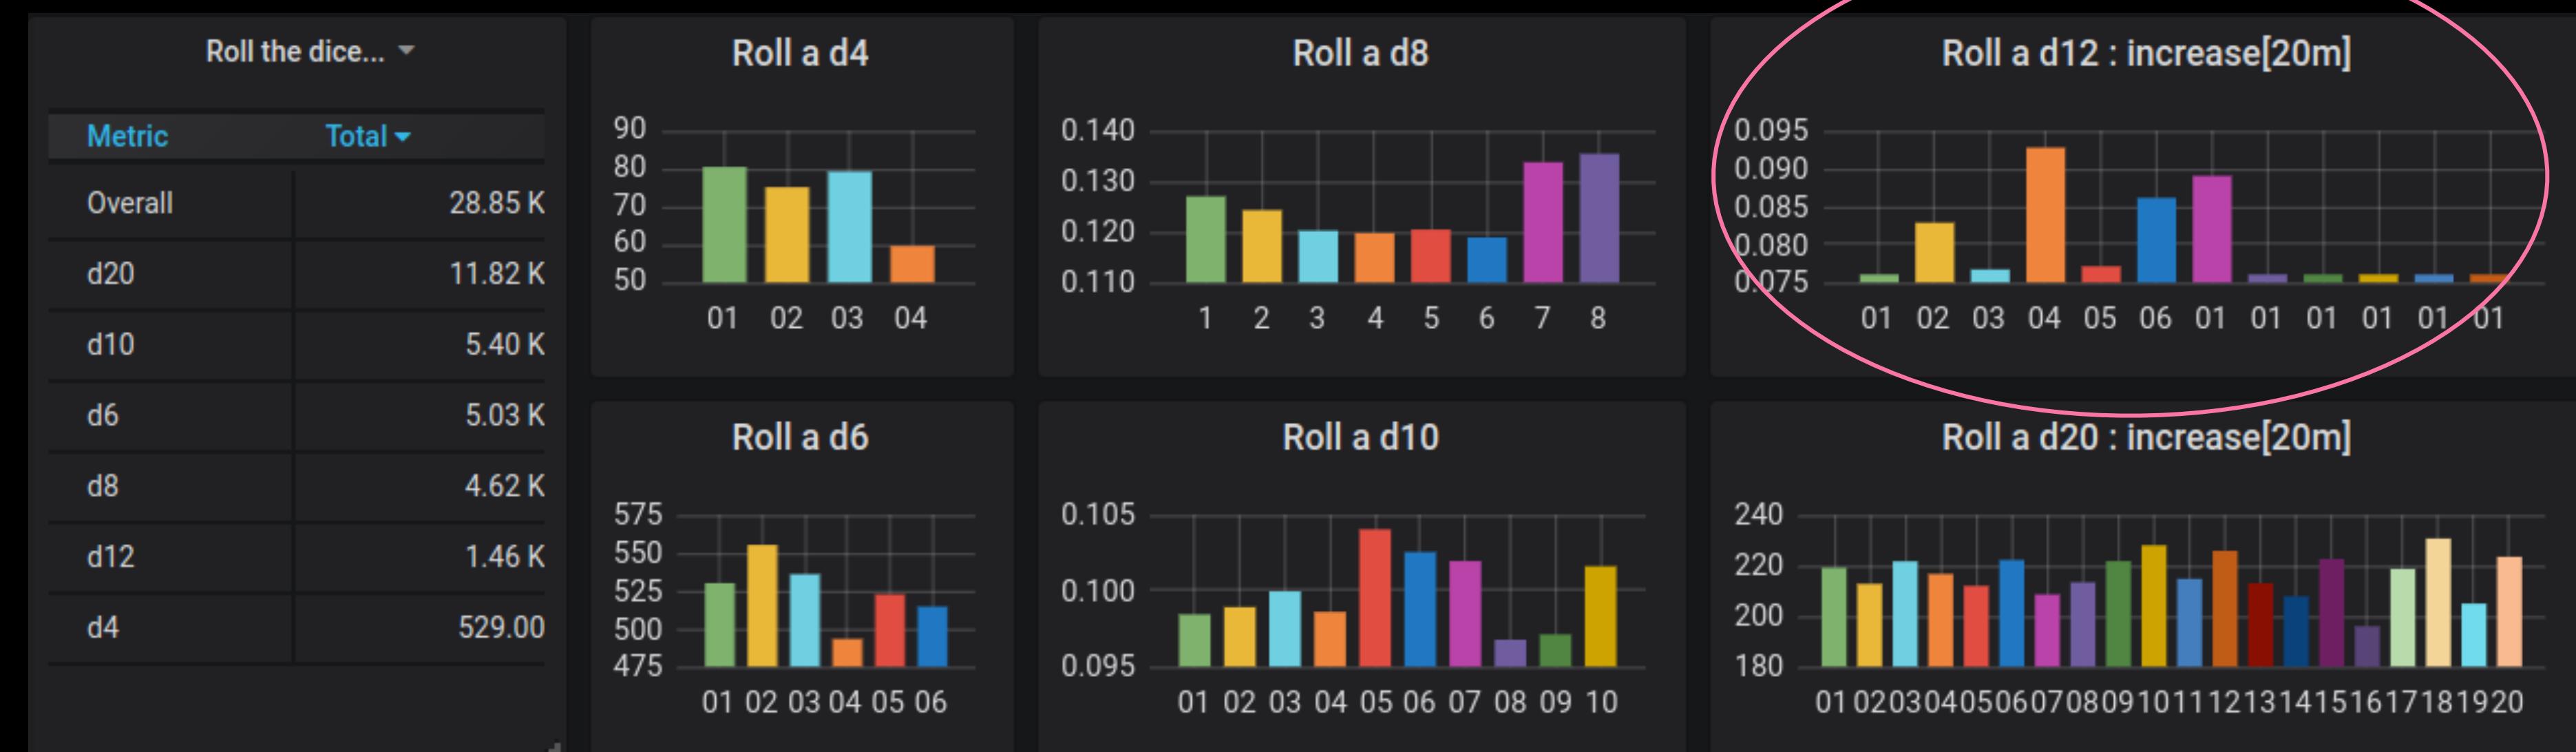


© Shutterstock / CiEll

Rolling dice

Micrometer:

```
Dice.setMonitor((k, v) -> {  
    registry.counter("dice.rolls", "die", k, "face", label(v))  
        .increment();  
});
```



Attacks : Hits and Misses

- **AC == Armor class**

Attacker rolls a d20...

20 – critical hit (HIT)

1 – critical miss (MISS)

- **DC == Difficulty class**

Defender rolls a d20...

: oneRound:

Troll(LARGE GIANT){**AC:15**,HP:84(8d10+40),STR:18(+4),DEX:13(+1),CON:20(+5),INT:7(-2),WIS:9(-1),C

Pit Fiend(LARGE FIEND){**AC:19**,HP:300(24d10+168),STR:26(+8),DEX:14(+2),CON:24(+7),INT:22(+6),V

: attack: **miss**: Troll(36) -> Pit Fiend(100)

: attack: **miss**: Troll(36) -> Pit Fiend(100)

: attack: **hit**> Troll(36) -> Pit Fiend(97) for 9 damage using Claws[7hit,11(2d6+4)|slashing]

: attack: **hit**> Pit Fiend(97) -> Troll(10) for 22 damage using Bite[14hit,22(4d6+8)|piercing]

: attack: **MISS**: Pit Fiend(97) -> Troll(10)

: attack: **HIT**> Pit Fiend(97) -> Troll(0) for 34 damage using Mace[14hit,15(2d6+8)|bludgeoning]

: oneRound: survivors

Pit Fiend(LARGE FIEND){AC:19,HP:300(24d10+168),STR:26(+8),DEX:14(+2),CON:24(+7),INT:22(+6),V

Micrometer: Timer & Distribution Summary

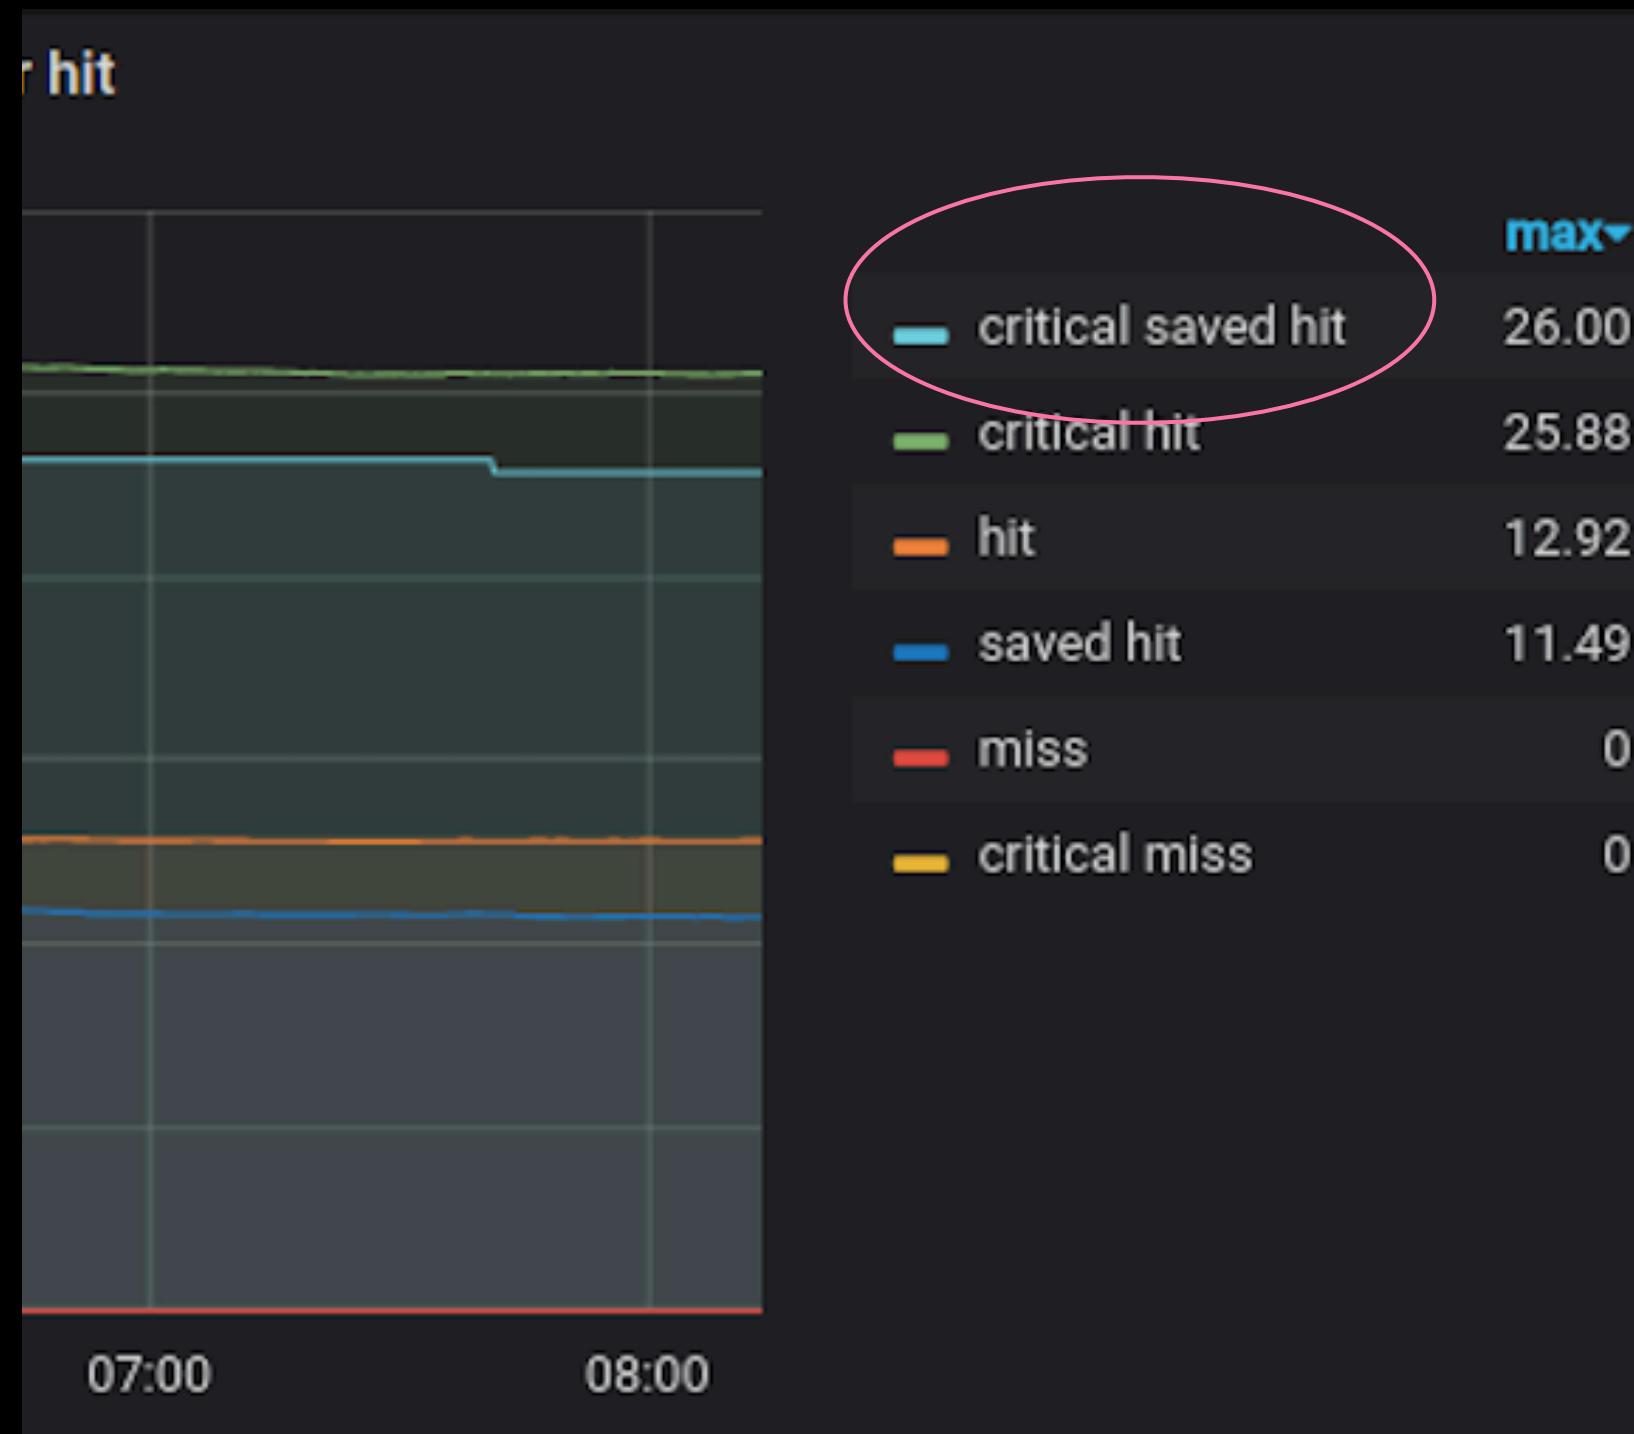
- Records value

```
registry.summary("round.attacks",
    "hitOrMiss", event.hitOrMiss(),
    "attackType", event.getAttackType(),
    "damageType", event.getType())
    .record((double) event.getDamageAmount());
```

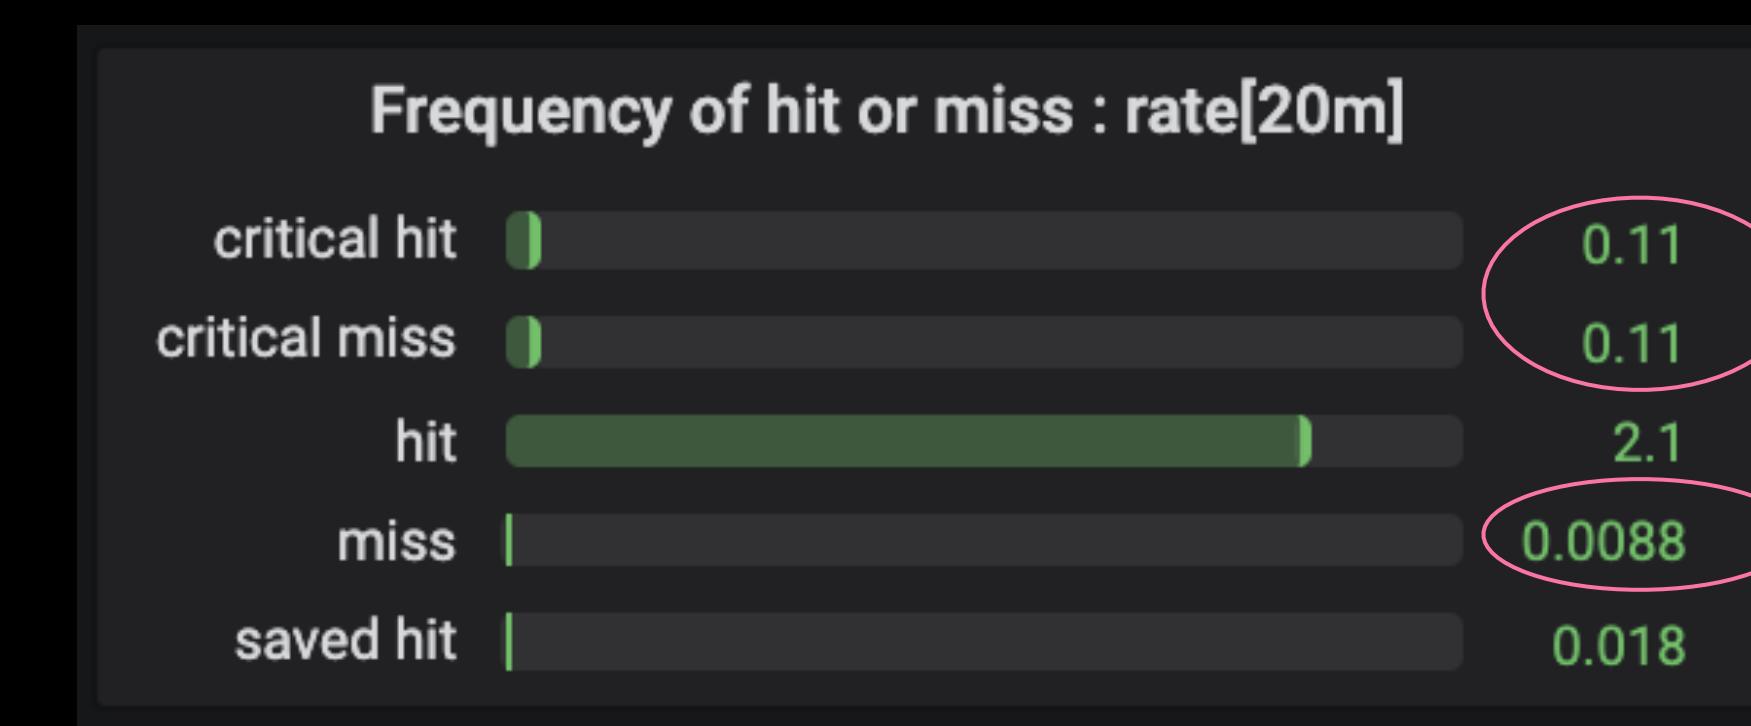
- Default metrics: count, sum, max
- Optional: cumulative histogram buckets, pre-computed quantile

So many bugs...

There is often no practical way to write tests for all permutations...



Who wrote this code, anyway?



This ended up being a problem with the original conversion from HTML to JSON.

Questions?